

# Project 1 SQL

The deadline for project 1 is:

**Fri 16 July, 5:00 pm (Sydney time)**

## 1. Aims

This project aims to give you practice in

- reading and understanding a moderately large relational schema (MyMyUNSW).
- implementing SQL queries and views to satisfy requests for information.
- The goal is to build some useful data access operations on the MyMyUNSW database. The data may contain some data inconsistencies; however, they won't affect your answers to the project.

## 2. How to do this project:

- Read this specification carefully and completely
- Familiarize yourself with the database **schema** (description, SQL schema, summary)
- Make a private directory for this project, and put a copy of the **proj1.sql** template there
- You **must** use the create statements in **proj1.sql** when defining your solutions
- Look at the expected outputs in the expected\_qX tables loaded as part of the **check.sql** file
- Solve each of the problems below, and put your completed solutions into **proj1.sql**
- Check that your solution is correct by verifying against the example outputs and by using the check\_qX() functions
- Test that your **proj1.sql** file will load *without error* into a database containing just the original MyMyUNSW data
- Double-check that your **proj1.sql** file loads in a *single pass* into a database containing just the original MyMyUNSW data
- Submit the project via give
- PLpgSQL functions are **not** allowed to use in this project
- For each question, you must output result within 120 seconds on Grieg server.

## 3. Introduction

All Universities require a significant information infrastructure in order to manage their affairs. This typically involves a large commercial DBMS installation. UNSW's student information system sits behind the MyUNSW web site. MyUNSW provides an interface to a PeopleSoft enterprise management system with an underlying Oracle database. This back-end system (Peoplesoft/Oracle) is often called NSS.

UNSW has spent a considerable amount of money (\$80M+) on the MyUNSW/NSS system, and it handles much of the educational administration plausibly well. Most people gripe about the quality of the MyUNSW interface, but the system does allow you to carry out most basic enrolment tasks online.

Despite its successes, MyUNSW/NSS still has a number of deficiencies, including:

- no waiting lists for course or class enrolment
- no representation for degree program structures
- poor integration with the UNSW Online Handbook

The first point is inconvenient, since it means that enrolment into a full course or class becomes a sequence of trial-and-error attempts, hoping that somebody has dropped out just before you attempt to enrol and that no-one else has grabbed the available spot.

The second point prevents MyUNSW/NSS from being used for three important operations that would be extremely helpful to students in managing their enrolment:

- finding out how far they have progressed through their degree program, and what remains to be completed
- checking what are their enrolment options for next semester (e.g., get a list of available courses)
- determining when they have completed all of the requirements of their degree program and are eligible to graduate

NSS contains data about student, courses, classes, pre-requisites, quotas, etc. but does not contain any representation of UNSW's degree program structures. Without such information in the NSS database, it is not possible to do any of the above three. So, in 2007 the COMP9311 class devised a data model that could represent program requirements and rules for UNSW degrees. This was built on top of an existing schema that represented all of the core NSS data (students, staff, courses, classes, etc.). The enhanced data model was named the MyMyUNSW schema.

The MyMyUNSW database includes information that encompasses the functionality of NSS, the UNSW Online Handbook, and the CATS (room allocation) database. The MyMyUNSW data model, schema and database are described in a separate document.

## 4. Setting Up

To install the MyMyUNSW database under your Grieg server, simply run the following two commands:

```
$ createdb proj1
$ psql proj1 -f /home/cs9311/web/21T2/proj/proj1/mymyunsw.dump
```

If you've already set up PLpgSQL in your template1 database, you will get one error message as the database starts to load:

```
psql:mymyunsw.dump:NN: ERROR: language "plpgsql" already exists
```

You can ignore this error message, but any other occurrence of ERROR during the load needs to be investigated.

If everything proceeds correctly, the load output should look something like:

```
SET
SET
SET
```

```

SET
SET
psql:mymyunsw.dump:NN: ERROR:  language "plpgsql" already exists
... if PLpgSQL is not already defined,
... the above ERROR will be replaced by CREATE LANGUAGE
SET
SET
SET
CREATE TABLE
CREATE TABLE
... a whole bunch of these
CREATE TABLE
ALTER TABLE
ALTER TABLE
... a whole bunch of these
ALTER TABLE

```

Apart from possible messages relating to plpgsql, you should get no error messages. The database loading should take less than 60 seconds on Grieg, assuming that Grieg is not under heavy load. (If you leave your project until the last minute, loading the database on Grieg will be considerably slower, thus delaying your work even more. The solution: at least load the database Right Now, even if you don't start using it for a while.) (Note that the mymyunsw.dump file is 50MB in size; copying it under your home directory or your /srvr directory is not a good idea).

If you have other large databases under your PostgreSQL server on Grieg or you have large files under your /srvr/YOU/ directory, it is possible that you will exhaust your Grieg disk quota. In particular, you will not be able to store two copies of the MyMyUNSW database under your Grieg server. The solution: remove any existing databases before loading your MyMyUNSW database.

If you are running PostgreSQL at home, you can download the files: mymyunsw.dump, proj1.sql to get you started. You can grab the check.sql separately, once it becomes available.

A useful thing to do initially is to get a feeling for what data is actually there. This may help you understand the schema better and will make the descriptions of the exercises easier to understand. Look at the schema. Ask some queries. Do it now.

Examples ...

```

$ psql proj1
... PostgreSQL welcome stuff ...
proj1=# \d
... look at the schema ...
proj1=# select * from Students;
... look at the Students table ...
proj1=# select p.unswid,p.name from People p join Students s on (p.id=s.id);
... look at the names and UNSW ids of all students ...
proj1=# select p.unswid,p.name,s.phone from People p join Staff s on (p.id=s.id);
... look at the names, staff ids, and phone #s of all staff ...
proj1=# select count(*) from Course_Enrolments;
... how many course enrolment records ...
proj1=# select * from dbpop();

```

... how many records in all tables ...  
proj1=# ... etc. etc. etc.  
proj1=# \q

You will find that some tables (e.g. Books, Requirements, etc.) are currently unpopulated; their contents are not needed for this project.

### Summary on Getting Started

To set up your database for this project, run the following commands in the order supplied:

```
$ createdb proj1
$ psql proj1 -f /home/cs9311/web/21T2/proj/proj1/mymyunsw.dump
$ psql proj1
... run some checks to make sure the database is ok
$ mkdir Project1Directory
... make a working directory for Project 1
$ cp /home/cs9311/web/21T2/proj/proj1/proj1.sql Project1Directory
```

The only error messages produced by these commands should be those noted above. If you omit any of the steps, then things will not work as planned.

### Notes

**Read these** before you start on the exercises:

- the marks reflect the relative difficulty/length of each question
- use the supplied **proj1.sql** template file for your work
- you may define as many additional functions and views as you need, provided that (a) the definitions in proj1.sql are preserved, (b) you follow the requirements in each question on what you are allowed to define
- make sure that your queries would work on any instance of the MyMyUNSW schema; don't customize them to work just on this database; we may test them on a different database instance
- do not assume that any query will return just a single result; even if it phrased as "most" or "biggest", there may be two or more equally "big" instances in the database
- when queries ask for people's names, use the `Person.name` field; it's there precisely to produce displayable names
- when queries ask for student ID, use the `People.unswid` field; the `People.id` field is an internal numeric key and of no interest to anyone outside the database
- unless specifically mentioned in the exercise, the order of tuples in the result does not matter; it can always be adjusted using `order by`. In fact, our `check.sql` will order your results automatically for comparison.
- the precise formatting of fields within a result tuple **does** matter; e.g. if you convert a number to a string using `to_char` it may no longer match a numeric field containing the same value, even though the two fields may look similar
- develop queries in stages; make sure that any sub-queries or sub-joins that you're using actually work correctly before using them in the query for the final view/function
- You can define SQL views to answer the following questions.

- If you meet with error saying something like “cannot change name of view column”, you can drop the view you just created by using command “**drop view VIEWNAME cascade;**” then create your new view again.

Each question is presented with a brief description of what's required. If you want the full details of the expected output, take a look at the expected\_qX tables supplied in the checking script.

## 5. Tasks

To facilitate the semi-auto marking, please pack all your SQL solutions into view as defined in each problem (see details from the solution template we provided).

### Q1 (3 marks)

Define a SQL view `Q1(subject_name)` gives the distinct names of subjects (refer to the `subjects.name`) whose subject code (refers to `subjects.code`) starts with ‘COMP6’.

### Q2 (3 marks)

Define a SQL view `Q2(room_name)` that gives the names of the rooms (refer to the `rooms.name`) that are ‘Meeting Room’ (refers to the `room_types.description`) in ‘Computer Science Building’ (refers to the `buildings.name`).

### Q3 (4 marks)

Define a SQL view `Q3(no, orgunit_name, program_num)` gives the names of Faculties and the number of distinct programs directly offered by each Faculty. The view should return the following details about each faculty:

- `no` is the incremental index number starting from 1. (See `row_number()` function in PostgreSQL Manual)
- `orgunit_name` should be taken from `orgunits.name` field.
- `program_num` as integer.

**Note:**

- The distinct programs are considered based on the `programs.id` field.
- The type of orgunit, i.e., Faculty, refers to the `orgunit_types.name`.
- The results should be ordered by `program_num` from lowest to highest.

### Q4 (4 marks)

Define a SQL view `Q4(course_id)` that gives the `id` of the courses in semester 2010 S1, which have 5 days of lecture classes per week. The view should return the following details about each course:

- `course_id` should be taken from `courses.id` field.

**Note:**

- The type of class, i.e., Lecture, refers to the `class_types.name`.

### Q5 (4 marks)

Define a SQL view `Q5(unswid, name)` that gives the `unswid` and `name` of all the distinct international students who got a mark of 100 (refers to the `course_enrolments.mark`) in any postgraduate course. And only consider the students whose `unswid` beginning with ‘31’. The view should return the following details about each student:

- `unswid` should be taken from `people.unswid` field.
- `name` should be taken from `people.name` field.

**Note:**

- An international student refers to the student having the type 'intl' (refers to `students.stype`).
- A postgraduate course refers to the subject of this course has the career 'PG' (refers to `subjects.career`).

#### Q6 (4 marks)

Define a SQL view `Q6(course_id, staff_name)` that gives the `id` of the courses whose classes are held at more than one campus in semester 2011 S1, and the name of all the `course_staff` of these courses. The view should return the following details:

- `course_id` should be taken from `courses.id` field.
- `staff_name` should be taken from `people.name` field.

**Note:**

- The results should not contain the courses with no `course_staff`.
- Only consider the courses offered by Faculty of Law.
- If a course has more than one course staff, return all of them.

#### Q7 (4 marks)

Define a SQL view `Q7(room_name, course_num)` that gives the name(s) of the room(s) with the highest `course_num` (refers to the number of courses held in that room). The view should return the following details about each room:

- `room_name` should be taken from `rooms.longname` field.
- `course_num` as integer.

**Note:**

- Only consider the 'Lecture Theatre' rooms, refers to `room_types.description`.
- If several rooms have the same `course_num` (which is highest), return all of them.

#### Q8(4 marks)

Define a SQL view `Q8(no, unswid)`, which gives the `unswid` of distinct international students enrolling in the years 2008 and 2009 (refers to `semesters.year`) in the COMPCS stream (refers to `streams.code`), but never enrolling in any course having 'GSOE' in its code (refers to `subjects.code`). The view should return the following details about each student:

- `no` is the incremental index number starting from 1.
- `unswid` should be taken from `people.unswid` field.

**Note:**

- Do not count duplicate records.
- The results should be ordered by `unswid` from highest to lowest.

#### Q9 (5 marks)

Database Systems course admins would like to know the maximum average marks among the semesters from year 2008 to 2009. Define a SQL view `Q9(year, term, max_avg_mark)` to help them find the maximum average mark among the above semesters. The view should return the following details:

- `year` of the `max_avg_mark`, should be taken from `semesters.year` field.
- `term` of the `max_avg_mark`, should be taken from `semesters.term` field.
- `max_avg_mark` as `numeric(4,2)`

Database Systems has value 'Database Systems' in the `subjects.name` field. You can find the

information about all the course offerings for a given subject from `courses`. You should calculate the average mark of enrolled students for a course offering from the table `course_enrolments`.

**Note:**

- There are two subjects that share the same name “Database Systems”, and we do not distinguish them in this question. In consequence, you may find more than one course for a single semester. In such case, there is no student enrolling in more than one course. Combine the courses with the same name into a single course for each semester before calculating the semester average mark.
- When calculating the average marks, only consider **not null** mark records.
- If there exist multiple courses with the same `max_avg_mark`, return all of them.

### Q10 (5 marks)

Below are the rules for students’ semester grade according to their semester average mark:

- i. HD if  $85 \leq \text{semester average mark}$ .
- ii. DN if  $75 \leq \text{semester average mark} < 85$ .
- iii. CR if  $65 \leq \text{semester average mark} < 75$ .
- iv. PS if  $50 \leq \text{semester average mark} < 65$ .
- v. FL otherwise.

Define SQL view `Q11(unswid, average_mark, grade)`, which gives the `unswid`, semester average mark and semester grade of the students with `unswid` beginning with ‘22’ in semester 2005 S1. The view should return the following details about each student:

- `unswid` should be taken from `people.unswid` field.
- `average_mark` should be calculated from their course marks in semester 2005 S1 as `numeric (4,2)`.
- `grade` should be defined as the rules given above. You will need to display ‘HD’, ‘DN’, ‘CR’, ‘PS’ or ‘FL’ for each student you found.

**Note:**

- For each student, we only consider the courses in which he/she receives a **not null** mark.
- Only consider the students enrolled in ‘Commerce’ program, refers to `programs.name`.
- FL only applies to the average mark rather than a single course mark in this case.
- Do not consider the students who never get any mark from any course in that semester.

### Q11 (5 marks)

Define SQL view `Q11(orgunit_name, intl_ratio)`, which gives the `unit_name` of 5 Faculties with highest ratio of enrolled distinct international students, and the corresponding ratio. If there are multiple faculties with the same `intl_ratio`, they have the same rank, which means your view may contain more than 5 rows. The view should return the following details about each unit:

- `orgunit_name` should be taken from `orgunits.longname` field.
- `intl_ratio` as `numeric (3,2)`.

Note that for this problem:

- You should ignore the units with no international student.
- `intl_ratio` should be calculated as the number of enrolled international students / the number of total enrolled students.
- We only count **program enrolments** when considering the number of student enrolments, and the programs should be offered by Faculties directly.
- Rank is with gaps. i.e., if there are 2 faculties ranking first, the third faculty will be ranked as third.

### Q12 (5 marks)

Define SQL view `Q12(unswid, name, course_num)`, which gives the `unswid` and `name` of the top 3 Associate Professors who has been a course lecturer with most times, and the number of courses. If multiple lecturers have the same number of courses, they have the same rank, which means your view may contain more than 3 rows. The view should return the following details about each theatre:

- `unswid` should be taken from `people.unswid` field.
- `name` should be taken from `people.name` field.
- `course_num` as integer

#### Note:

- A `people` as Associate Professor means his/her `staff_roles.name` is 'Associate Professor' in an affiliation.
- A `people` as course lecturer means he/she is in `course_staff` and `staff_roles.name` contains 'Lecturer'.
- Rank is with gaps. i.e., if there are 2 lecturers ranking first, the third lecturer will be ranked as third.

## 6. Submission

You can submit this project by doing the following:

- Students must submit an electronic copy of their answers to the above questions to the course website in Moodle.
- The file name should be `proj1_studentID.sql` (e.g., `proj1_z5100000.sql`).
- If you submit your project more than once, the last submission will replace the previous one
- In case that the system is not working properly, you must take the following actions:
- Please keep a copy of your submitted file on the CSE server. If you are not sure how, please have a look at [taggi](#).

The `proj1.sql` file should contain answers to all of the exercises for this project. It should be completely self-contained and able to load in a single pass, so that it can be auto-tested as follows:

- a fresh copy of the MyMyUNSW database will be created (using the schema from `mymyunsw.dump`)
- the data in this database may be **different** from the database that you're using for testing
- a new `check.sql` file may be loaded (with expected results appropriate for the database)
- the contents of your `proj1.sql` file will be loaded
- each checking function will be executed, and the results recorded

Before you submit your solution, you should check that it will load correctly for testing by using something like the following operations:

```
$ dropdb proj1          ... remove any existing DB
$ createdb proj1         ... create an empty database
$ psql proj1 -f /home/cs9311/web/21T2/proj/proj1/mymyunsw.dump ... load the MyMyUNSW
schema and data
$ psql proj1 -f /home/cs9311/web/21T2/proj/proj1/check.sql ... load the checking code
```



```
$ psql proj1 -f proj1.sql    ... load your solution
$ psql proj1
proj1=# select check_q1();    ... check your solution to question1
...
proj1=# select check_q6();    ... check your solution to question6
...
proj1=# select check_q12();   ... check your solution to question12
proj1=# select check_all();   ... check all your solutions
Note: if your database contains any views that are not available in a file somewhere, you should put
them into a file before you drop the database.
```

If your code loads with errors, fix it and repeat the above until it does not.

You must ensure that your proj1.sql file will load correctly (i.e. it has no syntax errors and it contains all of your view definitions in the correct order). If I need to manually fix problems with your proj1.sql file in order to test it (e.g. change the order of some definitions), you will be fined via half of the mark penalty for each problem. In addition, make sure that your queries are reasonably efficient. **For each question, you must output result within 120 seconds on Grieg server.** This time restriction applies to the execution of the 'select \* from check\_Qn()' calls. For each question, you will be fined via half of the mark penalty if your solution cannot output results within 120 seconds.

## 7. Late Submission Penalty

**20% reduction for each day late.**