

SQL

Assignment 1

Due: 5PM - Friday 25 June (Sydney Time)

Final Exam

- The university is asking for time and duration for final exam
- 2PM to 8PM Sydney time ?

SQL-99

SQL = Structured Query Language (pronounced “sequel”).

An ANSI/ISO standard language for querying and manipulating relational DBMSs.

Developed at IBM (San Jose Lab) during the 1970’s, and standardised during the 1980’s.

Appears that SQL will survive the rise of object-relational database systems.

Designed to be a “human readable” language supporting:

- relational algebra operations
- aggregation operations

Sample Database

To illustrate the features of SQL, we use a small example database below:

Beers(*name*, manf), Bars(*name*, addr, license)

Drinkers(*name*, addr, phone), Likes(*drinker, beer*)

Sells(*bar, beer*, price), Frequents(*drinker, bar*)

keys are in *italic* font and highlighted by underscore.

Sample Database_(cont)

Bars:

Name	Addr	License
Australia Hotel	The Rocks	123456
Coogee Bay Hotel	Coogee	966500
Lord Nelson	The Rocks	123888
Marble Bar	Sydney	122123
Regent Hotel	Kingsford	987654
Royal Hotel	Randwick	938500

Drinkers:

Name	Addr	Phone
Adam	Randwick	9385-4444
Gernot	Newtown	9415-3378
John	Clovelly	9665-1234
Justin	Mosman	9845-4321

Sample Database_(cont)

Beers:

Name	Manf
80/-	Caledonian
Bigfoot Barley Wine	Sierra Nevada
Burraborang Bock	George IV Inn
Crown Lager	Carlton
Fosters Lager	Carlton
Invalid Stout	Carlton
Melbourne Bitter	Carlton
New	Toohey's
Old	Toohey's
Old Admiral	Lord Nelson
Pale Ale	Sierra Nevada
Premium Lager	Cascade
Red	Toohey's
Sheaf Stout	Toohey's
Sparkling Ale	Cooper's
Stout	Cooper's
Three Sheets	Lord Nelson
Victoria Bitter	Carlton

Sample Database_(cont)

Frequents:

Drinker	Bar
Adam	Coogee Bay Hotel
Gernot	Lord Nelson
John	Coogee Bay Hotel
John	Lord Nelson
John	Australia Hotel
Justin	Regent Hotel
Justin	Marble Bar

Likes:

Drinker	Beer
Adam	Crown Lager
Adam	Fosters Lager
Adam	New
Gernot	Premium Lager
Gernot	Sparkling Ale
John	80/-
John	Bigfoot Barley Wine
John	Pale Ale
John	Three Sheets
Justin	Sparkling Ale
Justin	Victoria Bitter

Sample Database_(cont)

Sells:

Bar	Beer	Price
Australia Hotel	Burraborang Bock	3.5
Coogee Bay Hotel	New	2.25
Coogee Bay Hotel	Old	2.5
Coogee Bay Hotel	Sparkling Ale	2.8
Coogee Bay Hotel	Victoria Bitter	2.3
Lord Nelson	Three Sheets	3.75
Lord Nelson	Old Admiral	3.75
Marble Bar	New	2.8
Marble Bar	Old	2.8
Marble Bar	Victoria Bitter	2.8
Regent Hotel	New	2.2
Regent Hotel	Victoria Bitter	2.2
Royal Hotel	New	2.3
Royal Hotel	Old	2.3
Royal Hotel	Victoria Bitter	2.3

Example:

Beers:

Name	Manf
80/-	Caledonian
Bigfoot Barley Wine	Sierra Nevada
Burraborang Bock	George IV Inn
Crown Lager	Carlton
Fosters Lager	Carlton
Invalid Stout	Carlton
Melbourne Bitter	Carlton
New	Toohey's
Old	Toohey's
Old Admiral	Lord Nelson
Pale Ale	Sierra Nevada
Premium Lager	Cascade
Red	Toohey's
Sheaf Stout	Toohey's
Sparkling Ale	Cooper's
Stout	Cooper's
Three Sheets	Lord Nelson
Victoria Bitter	Carlton

SQL Queries: What beers are made by Toohey's?"

SELECT Name FROM Beers WHERE Manf = 'Toohey's';

SQL Queries

To answer the question “What beers are made by Toohey’s?”, we could ask:

```
SELECT Name FROM Beers WHERE Manf = 'Toohey's';
```

This gives a subset of the Beers relation, displayed as:

Name

New

Old

Red

Sheaf Stout

Quotes are escaped by doubling them (‘ ‘)

SQL Queries_(cont)

Query syntax is:

SELECT attributes

FROM relations

WHERE condition

The result of this statement is a table, which is typically displayed on output.

The SELECT statement contains the functionality of *select*, *project* and *join* from the relational algebra.

SQL Identifiers

Names are used to identify objects such as tables, attributes, views, ...

Identifiers in SQL use similar conventions to common programming languages:

- a sequence of alpha-numerics, starting with an alphabetic,
- not case-sensitive,
- reserve word disallowed, ...

SQL Keywords

Some of the frequently-used ones:

- ALTER AND CREATE
- FROM INSERT NOT OR
- SELECT TABLE WHERE

For PostgreSQL Keywords see the Appendix of PostgreSQL doc .

SQL Data Types

All attributes in SQL relations have domain specified.

SQL supports a small set of useful built-in data types: strings, numbers, dates, bit-strings.

Self defined data type is allowed in PostgreSQL.

Various type conversions are available:

- date to string, string to date, integer to real ...
- applied automatically “where they make sense”

SQL Data Types_(cont.)

Basic domain (type) checking is performed automatically.

Constraints can be used to “enforce” more complex domain membership conditions.

The NULL value is a member of all data types.

SQL Data Types_(cont.)

Comparison operators are defined on all types.

< > <= >= = !=

Boolean operators AND, OR, NOT are available within WHERE expressions to combine results of comparisons.

Comparison against NULL yields FALSE.

Can explicitly test for NULL using:

- *attr* IS NULL *attr* IS NOT NULL

Most data types also have type-specific operations available (e.g. arithmetic for numbers).

Which operations are actually applied depends on the implementation.

SQL Strings

Two kinds of string are available:

- `CHAR(n)` ... uses *n* bytes, left-justified, blank-padded
- `VARCHAR(n)` ... uses *0..n* bytes, no padding

String types can be coerced by blank-padding or truncation.

String literals are written using single quotes.

- `'John' = 'John ' != 'JOHN '`

String comparison

$str_1 < str_2$... compare using dictionary order

str LIKE $pattern$... matches string to pattern

Two kinds of pattern-matching:

- % matches anything (like *)
- _ matches any single char (like .)

Examples:

- | | |
|---------------------|----------------------------|
| ◦ Name LIKE 'Ja%' | Name begins with 'Ja' |
| ◦ Name LIKE '_i%' | Name has 'i' as 2nd letter |
| ◦ Name LIKE '%o%o%' | Name contains two 'o's |

String manipulation

string || *string* ... concatenate two strings

- 'Post' || 'greSQL' -> PostgreSQL

LENGTH(*str*) ... return length of string

SUBSTR(*str,start,length*) ... extract chars from within string

- substring('Thomas' from 2 for 3) -> hom

SQL Dates

Dates are simply specially-formatted strings, with a range of operations to implement date semantics.

Format is typically DD-Mon-YYYY, e.g. '18-Aug-1998'

Accepts other formats

Comparison operators implement before (<) and after (>).

(start1, end1) OVERLAPS (start2, end2)

- This expression yields true when two time periods (defined by their endpoints) overlap, false when they do not overlap.
- `SELECT (DATE '2001-02-16', DATE '2001-12-21') OVERLAPS (DATE '2001-10-30', DATE '2002-10-30');` -> *Result*: true

SQL Numbers

Various kinds of numbers are available:

smallint, int, bigint ... 2-bytes, 4-bytes and 8-bytes integers

real, double precision... 4-bytes and 8-bytes floating point

numeric(precision, scale)

- The *scale* of a numeric is the count of decimal digits in the fractional part, to the right of the decimal point.
- The *precision* of a numeric is the total count of significant digits in the whole number

SQL Numbers_(cont.)

Arithmetic operations:

- + - * / abs ceil floor power sqrt sin ...

Some operations apply to a column of numbers in a relation:

- *AVG(attr)* ... mean of values for *attr*
- *COUNT(attr)* ... number of rows in *attr* column
- *MIN/MAX(attr)* ... min/max of values for *attr*
- *SUM(attr)* ... sum of values for *attr*

Note: NULL value produces NULL result for arithmetic operation, but NULL is ignored in column operations.

Tuple and Set Literals

Tuple and set constants are both written as:

- (val1, val2, val3, ...)

The correct interpretation is worked out from the context.

Examples:

```
Student(stude#, name, course)
( 2177364, 'Jack Smith', 'BSc')    -- tuple literal
```

```
SELECT name
FROM Employees
WHERE job IN ('Lecturer', 'Tutor', 'Professor');    -- set literal
```


Querying a Single Relation

Formal semantics (relational algebra):

- start with relation R in FROM clause
- apply σ using Condition in WHERE clause
- apply π using Attributes in SELECT clause

SELECT *Attributes*

FROM R

WHERE *Conditions*

Querying a Single Relation_(cont.)

Operationally, we think in terms of a *tuple variable* ranging over all tuples of the relation.

Operational semantics:

FOR EACH tuple T in R DO

 check whether T satisfies the condition in the WHERE clause

 IF it does THEN

 print the attributes of T that are
 specified in the SELECT clause

 END

END

Projection by SQL

Assume a relation R and attributes $X \subseteq R$.

$\pi_X(R)$ is implemented in SQL as:

- SELECT X FROM R

Example:

Names of drinkers: $\pi_{Name}(Drinkers)$

- SELECT Name FROM Drinkers;

Name

Adam

Gernot

John

Justin

Drinkers:

Name	Addr	Phone
Adam	Randwick	9385-4444
Gernot	Newtown	9415-3378
John	Clovelly	9665-1234
Justin	Mosman	9845-4321

Projection by SQL_(cont.)

Example:

Names and addresses of drinkers = $\pi_{Name,Addr}(Drinkers)$

- SELECT Name, Addr FROM Drinkers;

NAME	ADDR
-----	-----
Adam	Randwick
Gernot	Newtown
John	Clovelly
Justin	Mosman

Projection by SQL_(cont.)

The symbol * denotes a list of all attributes.

Example:

All information about drinkers:

- SELECT * FROM Drinkers;

NAME	ADDR	PHONE
-----	-----	-----
Adam	Randwick	9385-4444
Gernot	Newtown	9415-3378
John	Clovelly	9665-1234
Justin	Mosman	9845-4321

Selection by SQL

$\sigma_{\text{Cond}}(\text{Rel})$ is implemented in SQL as:

SELECT * FROM Rel WHERE Cond

Example: Find the price that **Regent Hotel** charges for **New**

SELECT price

FROM Sells

WHERE bar = 'Regent Hotel' AND beer = 'New';

PRICE

2.2

The condition can be an arbitrarily complex boolean-valued expression using the operators mentioned previously.

Bar	Beer	Price
Australia Hotel	Burraborang Bock	3.5
Coogee Bay Hotel	New	2.25
Coogee Bay Hotel	Old	2.5
Coogee Bay Hotel	Sparkling Ale	2.8
Coogee Bay Hotel	Victoria Bitter	2.3
Lord Nelson	Three Sheets	3.75
Lord Nelson	Old Admiral	3.75
Marble Bar	New	2.8
Marble Bar	Old	2.8
Marble Bar	Victoria Bitter	2.8
Regent Hotel	New	2.2
Regent Hotel	Victoria Bitter	2.2
Royal Hotel	New	2.3
Royal Hotel	Old	2.3
Royal Hotel	Victoria Bitter	2.3

Selection by SQL_(cont.)

The “typical” SELECT query:

```
SELECT a1, a2, a3  
FROM Rel  
WHERE Cond
```

This corresponds to select followed by project:

$$\pi_{\{a1,a2,a3\}}(\sigma_{\text{Cond}}(Rel)).$$

Renaming via as

Ullman/Widom define a renaming operator ρ to avoid name clashes.

Example: $\rho_{Beers(Brand,Brewer)}(Beers)$

Gives a new relation, with same data as *Beers*, but with attribute names changed.

SQL provides *AS* to achieve this; it is used in the *SELECT* part.

Renaming via as_(cont.)

Example:

- Beers(name, manf)

```
SELECT name AS Brand, manf AS Brewer FROM Beers;
```

BRAND

80/-

Bigfoot Barley Wine

Burraborang Bock

Crown Lager

Fosters Lager

Invalid Stout

...

BREWER

Caledonian

Sierra Nevada

George IV Inn

Carlton

Carlton

Carlton

Expressions as Values in Columns

AS can also be used to introduce computed values

Example:

- Sells(bar, beer, price)

```
SELECT bar, beer, price*120 AS PriceInYen
FROM Sells;
```

BAR	BEER	PRICEINYEN
-----	-----	-----
Australia Hotel	Burraborang Bock	420
Coogee Bay Hotel	New	270
Coogee Bay Hotel	Old	300
Coogee Bay Hotel	Sparkling Ale	336
Coogee Bay Hotel	Victoria Bitter	276
...		

Just Display but no change to the database

Inserting Text in Result Table

Trick: to put text in output columns, use constant expression with *AS*.

Example:

Likes(drinker, beer)

```
SELECT drinker, 'likes Cooper's' AS WhoLikes
FROM Likes
WHERE beer = 'Sparkling Ale';
```

DRINKER	WHOLIKES
-----	-----
Gernot	likes Cooper's
Justin	likes Cooper's

Drinker	Beer
Adam	Crown Lager
Adam	Fosters Lager
Adam	New
Gernot	Premium Lager
Gernot	Sparkling Ale
John	80/-
John	Bigfoot Barley Wine
John	Pale Ale
John	Three Sheets
Justin	Sparkling Ale
Justin	Victoria Bitter

Find the brewers whose beers John likes.

```
SELECT Manf
FROM Likes, Beers
WHERE drinker = 'John' AND beer = name;
```

Likes:

Drinker	Beer
Adam	Crown Lager
Adam	Fosters Lager
Adam	New
Gernot	Premium Lager
Gernot	Sparkling Ale
John	80/-
John	Bigfoot Barley Wine
John	Pale Ale
John	Three Sheets
Justin	Sparkling Ale
Justin	Victoria Bitter

Beers:

Name	Manf
80/-	Caledonian
Bigfoot Barley Wine	Sierra Nevada
Burraborang Bock	George IV Inn
Crown Lager	Carlton
Fosters Lager	Carlton
Invalid Stout	Carlton
Melbourne Bitter	Carlton
New	Toohey's
Old	Toohey's
Old Admiral	Lord Nelson
Pale Ale	Sierra Nevada
Premium Lager	Cascade
Red	Toohey's
Sheaf Stout	Toohey's
Sparkling Ale	Cooper's
Stout	Cooper's
Three Sheets	Lord Nelson
Victoria Bitter	Carlton

Querying Multi-relations

Example: Find the brewers whose beers John likes.

- Likes(drinker, beer)
- Beers(name, manf)

```
SELECT Manf
FROM Likes, Beers
WHERE drinker = 'John' AND beer = name;
```

MANF

Caledonian
Sierra Nevada
Sierra Nevada
Lord Nelson

Note: could eliminate the duplicates by using *DISTINCT*.

Relational algebra: $\pi_{manf}(\sigma_{drinker='John'} Likes \bowtie_{(beer, name)} Beers)$.

Querying Multi-relations_(cont.)

Syntax:

SELECT *Attributes*

FROM *R1, R2, ...*

WHERE *Condition*

FROM clause contains a list of relations.

Querying Multi-relations_(cont.)

For SQL *SELECT* statement on several relations:

SELECT Attributes

FROM R1, R2, ...

WHERE Condition

Formal semantics (relational algebra):

- start with product $R1 \times R2 \times \dots$ in FROM clause
- apply σ using Condition in WHERE clause
- apply π using Attributes in SELECT clause

Querying Multi-relations_(cont.)

Operational semantics of *SELECT*:

```
FOR EACH tuple T1 in R1 DO
  FOR EACH tuple T2 in R2 DO
    ...
    check WHERE condition for current
    assignment of T1, T2, ... vars
    IF holds THEN
      print attributes of T1, T2, ...
      specified in SELECT      END
    END
  ...
END
```

For efficiency reasons, it is not implemented in this way!

Attribute Name Clashes

If a selection condition

- refers to two relations
- the relations have attributes with the same name

use the relation name to disambiguate.

Example: Which hotels have the same name as a beer?

```
SELECT Bars.name
```

```
FROM Bars, Beers
```

```
WHERE Bars.name = Beers.name;
```

Beers(name, manf)

Bars(name, addr, license)

None of them do, so the result is empty.

Attribute Name Clashes_(cont.)

Can use such qualified names, even if there is no ambiguity:

```
SELECT Sells.beer
```

```
FROM Sells
```

```
WHERE Sells.price > 3.00;
```

Advice:

- qualify attribute names only when absolutely necessary.
- SQL's AS operator cannot be used to resolve name clashes.

Table Name Clashes

The relation-dot-attribute convention doesn't help if we use the same relation twice in SELECT.

To handle this, we need to define new names for each “instance” of the relation in the FROM clause.

Example: Find pairs of beers by the same manufacturer.

Note: we should avoid:

- pairing a beer with itself e.g. (New,New)
- same pairs with different order e.g. (New,Old) (Old,New)

```
SELECT b1.name, b2.name
FROM Beers b1, Beers b2
WHERE b1.manf = b2.manf AND b1.name < b2.name;
```

NAME	NAME
-----	-----
Crown Lager	Fosters Lager
Crown Lager	Invalid Stout
Fosters Lager	Invalid Stout
Fosters Lager	Melbourne Bitter
....	

Beers:

Name	Manf
80/-	Caledonian
Bigfoot Barley Wine	Sierra Nevada
Burraborang Bock	George IV Inn
Crown Lager	Carlton
Fosters Lager	Carlton
Invalid Stout	Carlton
Melbourne Bitter	Carlton
New	Toohey's
Old	Toohey's
Old Admiral	Lord Nelson
Pale Ale	Sierra Nevada
Premium Lager	Cascade
Red	Toohey's
Sheaf Stout	Toohey's
Sparkling Ale	Cooper's
Stout	Cooper's
Three Sheets	Lord Nelson
Victoria Bitter	Carlton

Subqueries

The result of a SELECT-FROM-WHERE query can be used in the WHERE clause of another query.

Simplest Case: Subquery returns one tuple.

- Can treat the result as a constant value and use =.

Example: Find bars that sell New at the price same as the Coogee Bay Hotel charges for VB.

Sells:

Bar	Beer	Price
Australia Hotel	Burraborang Bock	3.5
Coogee Bay Hotel	New	2.25
Coogee Bay Hotel	Old	2.5
Coogee Bay Hotel	Sparkling Ale	2.8
Coogee Bay Hotel	Victoria Bitter	2.3
Lord Nelson	Three Sheets	3.75
Lord Nelson	Old Admiral	3.75
Marble Bar	New	2.8
Marble Bar	Old	2.8
Marble Bar	Victoria Bitter	2.8
Regent Hotel	New	2.2
Regent Hotel	Victoria Bitter	2.2
Royal Hotel	New	2.3
Royal Hotel	Old	2.3
Royal Hotel	Victoria Bitter	2.3



Subqueries_(cont.)

Example: Find bars that sell New at the price same as the Coogee Bay Hotel charges for VB.

```
SELECT bar
FROM Sells
WHERE beer = 'New'
      AND price =
          (SELECT price
           FROM Sells
           WHERE bar = 'Coogee Bay Hotel'
           AND beer = 'Victoria Bitter' );
```

BAR

Royal Hotel

Parentheses around the subquery are required.

NOT use subqueries

Example: Find bars that sell New at the price same as the Coogee Bay Hotel charges for VB.

```
SELECT b2.bar
FROM Sells b1, Sells b2
WHERE b1.beer = 'Victoria Bitter' and b1.bar = 'Coogee Bay Hotel' and
b1.price = b2.price and b2.beer = 'New';
```

BAR

Royal Hotel

Subqueries_(cont.)

Complex Case: Subquery returns multiple tuples/a relation.

- Treat it as a list of values, and use the various operators on lists/sets (e.g. IN).

IN Operator

Tests whether a specified tuple is contained in a relation.

tuple IN relation: is true iff the tuple is contained in the relation.

Conversely for *tuple* NOT IN relation.

Example: Find the name and brewers of beers that John likes.

Likes:

Drinker	Beer
Adam	Crown Lager
Adam	Fosters Lager
Adam	New
Gernot	Premium Lager
Gernot	Sparkling Ale
John	80/-
John	Bigfoot Barley Wine
John	Pale Ale
John	Three Sheets
Justin	Sparkling Ale
Justin	Victoria Bitter

Beers:

Name	Manf
80/-	Caledonian
Bigfoot Barley Wine	Sierra Nevada
Burraborang Bock	George IV Inn
Crown Lager	Carlton
Fosters Lager	Carlton
Invalid Stout	Carlton
Melbourne Bitter	Carlton
New	Toohey's
Old	Toohey's
Old Admiral	Lord Nelson
Pale Ale	Sierra Nevada
Premium Lager	Cascade
Red	Toohey's
Sheaf Stout	Toohey's
Sparkling Ale	Cooper's
Stout	Cooper's
Three Sheets	Lord Nelson
Victoria Bitter	Carlton

Subqueries_(cont.)

Example: Find the name and brewers of beers that John likes.

```
SELECT *  
FROM Beers  
WHERE name IN  
      (SELECT beer  
       FROM Likes  
       WHERE drinker = 'John'  
      );
```

NAME -----	MANF -----
80/-	Caledonian
Bigfoot Barley Wine	Sierra Nevada
Pale Ale	Sierra Nevada
Three Sheets	Lord Nelson

- The subquery answers the question "What are the names of the beers that John likes?"
- Note that this query can be answered equally well without using IN.
- The subquery version is potentially (but not always) less efficient.

Subqueries_(cont.)

Example: Find the name and brewers of beers that John likes.

```
SELECT *  
FROM Beers  
WHERE name IN  
      (SELECT beer  
       FROM Likes  
       WHERE drinker = 'John'  
      );
```

```
SELECT Beers.*  
FROM Beers, Likes  
Where Beers.name = Likes.beer and  
Likes.drinker = 'John';
```

NAME	MANF
-----	-----
80/-	Caledonian
Bigfoot Barley Wine	Sierra Nevada
Pale Ale	Sierra Nevada
Three Sheets	Lord Nelson

Example: Find the beers uniquely made by their manufacturer.

Beers:

Name	Manf
80/-	Caledonian
Bigfoot Barley Wine	Sierra Nevada
Burraborang Bock	George IV Inn
Crown Lager	Carlton
Fosters Lager	Carlton
Invalid Stout	Carlton
Melbourne Bitter	Carlton
New	Toohey's
Old	Toohey's
Old Admiral	Lord Nelson
Pale Ale	Sierra Nevada
Premium Lager	Cascade
Red	Toohey's
Sheaf Stout	Toohey's
Sparkling Ale	Cooper's
Stout	Cooper's
Three Sheets	Lord Nelson
Victoria Bitter	Carlton

EXISTS Function

EXISTS(relation) is true iff the relation is non-empty.

Example: Find the beers uniquely made by their manufacturer.

```
SELECT name
FROM Beers b1
WHERE NOT EXISTS
      (SELECT *
       FROM Beers
       WHERE manf = b1.manf
       AND name != b1.name
      );
```

NAME

80/-

Burraborang Bock

Premium Lager

A subquery that refers to values from a surrounding query is called a *correlated subquery*.

Quantifiers

ANY and ALL behave as existential and universal quantifiers respectively.

Example: Find the beers sold for the highest price.

```
SELECT beer
FROM Sells
WHERE price >=
    ALL(
        SELECT price
        FROM sells
    );
```

BEER

Three Sheets

Old Admiral

Beware: in common use, "any" and "all" are often synonyms.

E.g. "I'm better than any of you" vs. "I'm better than all of you".

Find the drinkers and beers such that the drinker likes the beer and frequents a bar that sells it.

Sells

Bar	Beer	Price
Australia Hotel	Burraborang Bock	3.5
Coogee Bay Hotel	New	2.25
Coogee Bay Hotel	Old	2.5
Coogee Bay Hotel	Sparkling Ale	2.8
Coogee Bay Hotel	Victoria Bitter	2.3
Lord Nelson	Three Sheets	3.75
Lord Nelson	Old Admiral	3.75
Marble Bar	New	2.8
Marble Bar	Old	2.8
Marble Bar	Victoria Bitter	2.8
Regent Hotel	New	2.2
Regent Hotel	Victoria Bitter	2.2
Royal Hotel	New	2.3
Royal Hotel	Old	2.3
Royal Hotel	Victoria Bitter	2.3

Likes

Drinker	Beer
Adam	Crown Lager
Adam	Fosters Lager
Adam	New
Gernot	Premium Lager
Gernot	Sparkling Ale
John	80/-
John	Bigfoot Barley Wine
John	Pale Ale
John	Three Sheets
Justin	Sparkling Ale
Justin	Victoria Bitter

Frequents

Drinker	Bar
Adam	Coogee Bay Hotel
Gernot	Lord Nelson
John	Coogee Bay Hotel
John	Lord Nelson
John	Australia Hotel
Justin	Regent Hotel
Justin	Marble Bar

Union, Intersection, Difference

R1 UNION R2: produces the union of the two relations R1 and R2.

Similarly for R1 INTERSECT R2 and R1 Except R2.

Example: Find the drinkers and beers such that the drinker likes the beer and frequents a bar that sells it.

```
(SELECT *  
  FROM Likes  
)  
INTERSECT  
(SELECT drinker,beer  
  FROM Sells, Frequents  
  WHERE Frequents.bar = Sells.bar  
);
```

DRINKER	BEER
-----	-----
Adam	New
John	Three Sheets
Justin	Victoria Bitter

Divide Operation

Find bars each of which sell all beers Justin likes.

Relational Algebra: $\pi_{bar,beer} Sells \div (\pi_{beer}(\sigma_{drinker='Justin'} Likes))$

Bar	Beer	Price
Australia Hotel	Burraborang Bock	3.5
Coogee Bay Hotel	New	2.25
Coogee Bay Hotel	Old	2.5
Coogee Bay Hotel	Sparkling Ale	2.8
Coogee Bay Hotel	Victoria Bitter	2.3
Lord Nelson	Three Sheets	3.75
Lord Nelson	Old Admiral	3.75
Marble Bar	New	2.8
Marble Bar	Old	2.8
Marble Bar	Victoria Bitter	2.8
Regent Hotel	New	2.2
Regent Hotel	Victoria Bitter	2.2
Royal Hotel	New	2.3
Royal Hotel	Old	2.3
Royal Hotel	Victoria Bitter	2.3

Drinker	Beer
Adam	Crown Lager
Adam	Fosters Lager
Adam	New
Gernot	Premium Lager
Gernot	Sparkling Ale
John	80/-
John	Bigfoot Barley Wine
John	Pale Ale
John	Three Sheets
Justin	Sparkling Ale
Justin	Victoria Bitter

Divide Operation

Find bars each of which sell all beers Justin likes.

Relational Algebra: ~~$Sells \div (\pi_{beer}(\sigma_{drinker='Justin'} Likes))$~~

$$\pi_{bar,beer} Sells \div (\pi_{beer}(\sigma_{drinker='Justin'} Likes))$$

select distinct a.bar

from sells a

where not exists

((select b.beer from likes b
where b.drinker = 'Justin')

except

(select c.beer from sells c
where c.bar = a.bar)

);

BAR

Coogee Bay Hotel

Aggregation

Selection clauses can contain aggregation operations.

Example: What is the average price of New?

```
SELECT AVG(price)  ← AVG (DISTINCT price)
FROM Sells
WHERE beer = 'New';
```

```
AVG(PRICE)
```

```
-----
```

```
2.3875
```

All prices for 'New' will be included, even if two hotels sell it at the same price.

If set semantics used, the result would be wrong.

Aggregation_(cont.)

If we want set semantics, we can force using DISTINCT.

Example: How many different bars sell beer?

```
SELECT COUNT(DISTINCT bar)
FROM Sells;
```

```
COUNT(DISTINCTBAR)
```

```
-----
```

```
6
```

Without DISTINCT, the result is 15 ... the number of entries in the Sells table.

Aggregation_(cont.)

The following operators apply to a list of numeric values in one column of a relation:

- SUM AVG MIN MAX COUNT

The notation COUNT(*) gives the number of tuples in a relation.

Example: How many different beers are there?

```
SELECT COUNT(*) FROM Beers;
```

```
COUNT(*)
```

```
-----
```

```
18
```

Grouping

SELECT-FROM-WHERE can be followed by *GROUP BY* to:

- partition result relation into groups (according to values of specified attribute)
- treat each group separately in computing aggregations

Example: How many beers does each brewer make?

```
SELECT manf, COUNT(beer)
FROM Beers
GROUP BY manf;
```

MANF	COUNT(beer)
-----	-----
Caledonian	1
Carlton	5
Cascade	1
Cooper's	2
George IV Inn	1
Lord Nelson	2
Sierra Nevada	2
Toohey's	4

Grouping_(cont.)

GROUP BY is used as follows:

SELECT *attributes/aggregations*

FROM *relations*

WHERE *condition*

GROUP BY *attribute*

Semantics:

- partition result into groups based on distinct values of attribute
- apply any aggregation separately to each group

Grouping_(cont.)

Grouping is typically used in queries involving the phrase “for each”.

Example: For each drinker, find the average price of New at the bars they frequently go to.

```
SELECT drinker, AVG(price)
FROM Frequent, Sells
WHERE beer = 'New' AND Frequent.bar = Sells.bar
GROUP BY drinker;
```

DRINKER	AVG(PRICE)
-----	-----
Adam	2.25
John	2.25
Justin	2.5

Grouping_(cont.)

When using grouping, every attribute in the SELECT list must:

- have an aggregation operator applied to it OR
- appear in a GROUP-BY clause

Incorrect Example: Find the cheapest beer price in each bar.

```
SELECT bar, MIN(price)
```

```
FROM Sells;
```

ERROR: column "sells.bar" must appear in the GROUP BY clause or be used in an aggregate function

LINE 1: select bar, min(price) from sells;

Grouping_(cont.)

How to answer the above query?

```
SELECT bar, MIN(price)
FROM Sells
GROUP BY BAR
```

bar	MIN(PRICE)
-----	-----
Australia Hotel	3.5
Coogee Bay Hotel	2.25
Lord Nelson	3.75
Marble Bar	2.8
Regent Hotel	2.2
Royal Hotel	2.3

Eliminating Groups

In some queries, you can use the WHERE condition to eliminate groups.

Example: Average beer price by suburb excluding hotels in The Rocks.

```
SELECT Bars.addr, AVG(Sells.price)
FROM Sells, Bars
WHERE Bars.addr != 'The Rocks'
AND Sells.bar = Bars.name
GROUP BY Bars.addr;
```

ADDR	AVG(SELLS.PRICE)
-----	-----
Coogee	2.4625
Kingsford	2.2
Randwick	2.3
Sydney	2.8

Eliminating Groups_(cont.)

For more complex conditions on groups, use the HAVING clause.

HAVING is used to qualify a GROUP-BY clause:

```
SELECT attributes/aggregations  
FROM relations  
WHERE condition (on tuples)  
GROUP BY attribute  
HAVING condition (on group);
```

Semantics of HAVING:

- generate the groups as for GROUP-BY
- eliminate any group not satisfying HAVING condition
- apply an aggregation to remaining groups

Eliminating Groups_(cont.)

Example: Find the average price of popular beers (i.e. those that are served in more than one hotel).

```
SELECT beer, AVG(price)
FROM Sells
GROUP BY beer
HAVING COUNT(bar) > 1;
```

BEER	AVG(PRICE)
-----	-----
New	2.3875
Old	2.533333333
Victoria Bitter	2.4

Defining a Database Schema

Relations (tables) are created using:

```
CREATE TABLE RelName (  
    attribute1 ~ domain1 ~ properties  
    attribute2 ~ domain2 ~ properties  
    attribute3 ~ domain3 ~ properties  
    ...  
)
```

where properties can include details about primary keys,
foreign keys, default values, and constraints on attribute values.

Tables are removed via **DROP TABLE** *RelName*;

Defining a Database Schema_(cont.)

Example:

```
CREATE TABLE Beers (  
    name VARCHAR(20) PRIMARY KEY,  
    manf VARCHAR(20),  
);  
  
CREATE TABLE Bars (  
    name VARCHAR(30) PRIMARY KEY,  
    addr VARCHAR(30),  
    license INTEGER  
);
```


Declaring Keys

Primary keys:

- if a single attribute, declare with attribute
- if several attributes, declare at end of attribute list

For attributes which have distinct values for each tuple, can note this via:

- *attribute domain* UNIQUE

Declaring Keys_(cont.)

Declaring foreign keys assures referential integrity.

Foreign a key:

- specify Relation (Attribute) to which it refers.

For instance, if we want to delete a tuple from Beers, and there are tuples in Sells that refer to it, we could either:

- **reject** the deletion
- **cascade** the deletion and remove Sells records
- **set-NULL** the foreign key attribute

Can force cascade via *ON DELETE CASCADE* after *REFERENCES*.

Other Attribute Properties

Can specify that an attribute is not allowed to be *NULL*.

This property applies automatically to *PRIMARY KEY* attributes.

Can specify a *DEFAULT* value which will be assigned if none is supplied during insert.

Example:

```
CREATE TABLE Likes (  
    drinker VARCHAR(20) DEFAULT 'Joe',  
    beer VARCHAR(30) DEFAULT 'New',  
    PRIMARY KEY(drinker, beer)  
);
```

Other Attribute Properties_(cont.)

In fact, *NOT NULL* is a special case of a constraint on the value that an attribute is allowed to take.

SQL has a more general mechanism for specifying such constraints.

- *attr_name type CHECK (condition)*

The Condition can be arbitrarily complex, and may even involve other attributes, relations and *SELECT* queries.

Other Attribute Properties_(cont.)

Example:

```
CREATE TABLE Example
```

```
(
```

```
    gender CHAR(1) CHECK (gender IN ('M','F')),
```

```
    Xvalue INT NOT NULL,
```

```
    Yvalue INT CHECK (Yvalue > Xvalue),
```

```
    Zvalue FLOAT CHECK (Zvalue > ( SELECT MAX(price)
```

```
                                FROM Sells))
```

```
);
```

Database Modification

Simple Insertion

Accomplished via the INSERT operation:

```
INSERT INTO Relation VALUES
```

```
(val1, val2, val3, ...)
```

Example: Add the fact that Justin likes 'Old'.

```
INSERT INTO Likes VALUES ('Justin', 'Old');
```

Can re-order attributes in tuple constant as long as order is specified in the INTO clause.

```
INSERT INTO Sells(price,bar,beer) VALUES
```

```
(2.50, 'Coogee Bay Hotel', 'Pale Ale');
```

Simple Insertion

Example: insertion with insufficient values.

E.g. we specify that drinkers' phone numbers cannot be NULL.

```
ALTER TABLE Drinkers ALTER COLUMN phone SET NOT NULL;
```

And then try to insert a new drinker whose phone number we don't know:

```
INSERT INTO Drinkers(name,addr)
```

```
VALUES ('Zoe', 'Manly');
```

ERROR: null value in column "phone" violates not-null constraint

DETAIL: Failing row contains (Zoe, Manly, null).

Insertion from Queries

Can use the result of a query to perform insertion of multiple tuples at once.

```
INSERT INTO Relation ( Subquery );
```

Tuples of Subquery must be projected into a suitable format (i.e. matching the tuple-type of Relation).

Insertion from Queries_(cont.)

Example: Create a relation of John's potential drinking buddies (i.e. people who go to the same bars as John).

```
CREATE TABLE DrinkingBuddies (  
    name varchar(20)  
);  
  
INSERT INTO DrinkingBuddies  
(  
    SELECT DISTINCT f2.drinker  
    FROM Frequents f1, Frequents f2  
    WHERE f1.drinker = 'John'  
        AND f2.drinker != 'John'  
        AND f1.bar = f2.bar  
);
```

Deletion

Accomplished via the DELETE operation:

DELETE FROM Relation

WHERE *Condition*

Removes all tuples from Relation that satisfy Condition.

Example: Justin no longer likes Sparkling Ale.

DELETE FROM Likes

WHERE drinker = 'Justin'

AND beer = 'Sparkling Ale';

Special case: Make relation R empty.

DELETE FROM R;

Deletion_(cont.)

Example: Delete all beers for which there is another beer by the same manufacturer.

```
DELETE FROM Beers b
WHERE EXISTS
  ( SELECT name
    FROM Beers
    WHERE manf = b.manf
    AND name != b.name);
```

Semantics here is subtle ...

If there is a manufacturer that makes only two beers, how many of them will be deleted?

E.g. after first beer is deleted, second beer no longer satisfies condition.

In fact, condition is evaluated for each tuple before making any changes.

Deletion_(cont.)

Semantics of the above Deletion:

Evaluation of DELETE FROM R WHERE Cond can be viewed as:

```
FOR EACH tuple T in R DO
    IF T satisfies Cond THEN
        make a note of this T
    END
END
FOR EACH noted tuple T DO
    remove T from relation R
END
```

Updates

An update allows you to modify values of specified attributes in specified tuples of a relation:

UPDATE *R*

SET *list of assignments*

WHERE *Condition*

Each tuple in relation *R* that satisfies *Condition* has the assignments applied to it.

Example: John moves to Coogee.

UPDATE Drinkers

SET addr = 'Coogee' ,

 phone = '9665-4321'

WHERE name = 'John';

Updates_(cont.)

Can update many tuples at once (all tuples that satisfy condition)

“Good” Example: Make \$3 the maximum price for beer.

```
UPDATE Sells
```

```
SET price = 3.00
```

```
WHERE price > 3.00;
```

“Bad” Example: Increase beer prices by 10%.

```
UPDATE Sells
```

```
SET price = price * 1.10;
```

Changing Tables

Accomplished via the ALTER TABLE operation:

- ALTER TABLE *Relation Modifications*

Some possible modifications are:

- add a new column (attribute),
- change the properties of an existing attribute,
- remove an attribute

Changing Tables_(cont.)

Example: Add phone numbers for hotels.

```
ALTER TABLE Bars
```

```
ADD phone char(10) DEFAULT 'Unlisted';
```

This appends a new column to the table and sets value for this attribute to 'Unlisted' in every tuple.

Specific phone numbers can subsequently be added via:

```
UPDATE Bars
```

```
SET phone = '9665-0000'
```

```
WHERE name = 'Coogee Bay Hotel';
```

If no default values is given, new column is set to all NULL.

Changing Tables_(cont.)

Can make multiple changes to one relation with a single ALTER.

Example: Add opening and closing times to Bars

```
ALTER TABLE Bars
```

```
Add opens NUMERIC(4,2) DEFAULT 10.00 ,
```

```
Add closes NUMERIC(4,2) DEFAULT 23.00 ,
```

```
Add manager VARCHAR(20)
```

```
;
```

Note that manager will be initially *NULL* for all hotels.

Views

A **view** is like a "virtual relation" defined in terms of other relations.

The other relations may be views (*intensional relations*) or stored relations (*extensional relations, base relations*).

View are defined via: `CREATE VIEW ViewName AS Query`

The view is valid only as long as the underlying query is valid.

Views may be removed via: `DROP VIEW ViewName`

Removing a view has no effect on the relations used by the view.

Views_(cont.)

Example: An avid CUB drinker might not be interested in any other kinds of beer.

```
CREATE VIEW MyBeers AS
  SELECT name, manf
  FROM Beers
  WHERE manf = 'Carlton';
SELECT * FROM MyBeers;
```

NAME	MANF
-----	-----
Crown Lager	Carlton
Fosters Lager	Carlton
Invalid Stout	Carlton
Melbourne Bitter	Carlton
Victoria Bitter	Carlton

Views_(cont.)

A view might not use all attributes of the base relations.

Example: We don't really need the address of inner-city hotels.

```
CREATE VIEW InnerCityHotels AS
  SELECT name, license
  FROM Bars
  WHERE addr = 'The Rocks' OR addr = 'Sydney';
SELECT * FROM InnerCityHotels;
```

NAME	LICENSE
-----	-----
Australia Hotel	123456
Lord Nelson	123888
Marble Bar	122123

Renaming View Attributes

This can be achieved in two different ways:

```
CREATE VIEW InnerCityPubs AS

    SELECT name AS pub, license AS lic

    FROM Bars

    WHERE addr IN ('The Rocks', 'Sydney');

CREATE VIEW InnerCityPubs(pub,lic) AS

    SELECT name, license

    FROM Bars

    WHERE addr IN ('The Rocks', 'Sydney');
```

Querying Views

Views can be used in queries just as if they were stored relations.

Unlike stored relations, views can "change" without explicit modification operations (i.e. by changing underlying relations).

Example: The Lord Nelson changes license.

```
UPDATE Bars SET license='111223' WHERE name='Lord Nelson'  
SELECT * FROM InnerCityHotels;
```

NAME	LICENSE
-----	-----
Australia Hotel	123456
Marble Bar	12212
Lord Nelson	111223

Querying Views_(cont.)

We can treat views as "macros" that will be re-written into queries on the base relation.

This is most easily seen by converting to relational algebra, and following transformation that an SQL query evaluator might make.

Example: Using the InnerCityHotels view.

```
CREATE VIEW InnerCityHotels AS
    SELECT name, license
    FROM Bars
    WHERE addr IN ('The Rocks', 'Sydney');
SELECT pub FROM InnerCityHotels WHERE lic = '123456';
```

Updating Views

Under the following conditions, it makes sense to allow view updates:

- the view involves a single relation R
- the WHERE clause does not involve R in a subquery
- there must be attributes in SELECT that allow the new tuple to be retrieved; unmentioned attributes are set to NULL

Updating Views_(cont.)

Example: Our InnerCityHotel view is not updatable.

```
INSERT INTO InnerCityHotels
```

```
VALUES ('Jackson''s on George', '9876543');
```

creates a new tuple in the Bars relation:

```
('Jackson''s on George', NULL, '9876543')
```

when we SELECT from the view, this new tuple does not satisfy the view condition:

```
addr IN ('The Rocks', 'Sydney')
```

Updating Views_(cont.)

If we had chosen to omit the license attribute instead, it would be updatable:

```
CREATE VIEW CityHotels AS
  SELECT name,addr FROM Bars
  WHERE addr IN ('The Rocks', 'Sydney');
INSERT INTO CityHotels
  VALUES ('Jackson''s on George', 'Sydney');
SELECT * FROM CityHotels;
```

NAME	ADDR
-----	-----
Australia Hotel	The Rocks
Marble Bar	Sydney
Jackson's on George	Sydney

Learning Outcomes

- Basic queries in SQL
- Querying multiple tables
- Aggregate queries
- Insert/delete/update
- Change schemas
- Views